

UrbanPark: Street Parking Detection through Computer Vision

Cynthia Chen , Abhinav Agarwal

Stanford University

cchen24@stanford.edu, abhinav4@stanford.edu

1 Introduction

In the field of object detection, one common method is parking space detection and classification. Upon doing research, we found several online datasets relating to parking space detection, several of which are on Kaggle. These datasets include images of original parking lots, boxes - which are the bounding box labels for the original images, and finally an annotations.xml file, which contain the coordinates of the bounding boxes and labels. Typically, work has been done in terms of training a CV model on these images and predicting the annotations and bounding boxes for the images. The previous research/data that has been done uses pytorch via a R-CNN model to predict bounding boxes for parking spots, and indicate if they are red or green depending on if there is an open or closed parking space. This method has been found to be very effective so far, in some applications. However, my partner and I found that there was a lack of research done on street parking classification and identification, so we have decided to make that the focus of our final project.

2 Background and Related Work

The first thing that we did was to find the datasets that could work for our street parking detection problem. After looking through a few datasets on Kaggle and other websites, we encountered one on Github that contained a large repository of 450+ images. (We used the starter code from this github repository to print the images and set up the baseline RCNN model). What we discovered upon looking at these images was that it contained a mix of both parking lot images and street images, which was good because we were initially looking to create street images in the first place, and after discovering these we just needed to change the angle of them so that they would look more similar to the parking lot datasets. As such, related work is mostly centered on tackling parking lot datasets but not street parking datasets. Additionally, one of our goals was to make the RCNN model more robust and see whether adding more complexity to the model would help it to garner better results, which is something that other works have not done in great depth. Finally, we have not seen past works perform perspective warping on this type of problem, so we aim to introduce this as a layer of novelty.

First, we created a Google Colaboratory notebook to load the training and validation datasets for the parking and street im-



Figure 1: Demonstrative Example of Bounding Boxes

ages. We were able to separate the data into a training dataset, validation dataset, and test dataset. Each image has three components: an image batch, rois batch, and labels batch. (We were able to set this up using the starter code that the Github repository provided, and load them into three variables per image.) The rois batch represented the red and green bounding boxes that would be overlaid on top of the cars in the image, and the labels represented whether those spaces/bounding boxes represented empty or occupied parking spots.

3 UrbanPark

3.1 Baseline RCNN Model

We used the RCNN model from the github repository as our baseline for evaluation. However, we wrote a function to compute the accuracy of the model on the testing dataset by taking the percent of bounding boxes that differed from the ground truth bounding boxes of the images.

Overall, this baseline RCNN model leverages a pre-trained ResNet-50 backbone for feature extraction, customizes ROI pooling, and allows fine-tuning of specific layers to adapt the model for this specific task. The ResNet-50 backbone is used as a feature extractor, and the model is initialized with pre-trained weights obtained from training on the large-scale dataset. The last fully connected layer (fc layer) of the ResNet-50 is modified to output 2 classes, for

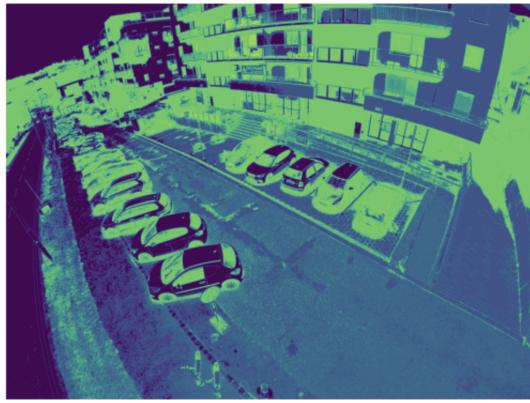


Figure 2: K-Means for Four Different Features



Figure 3: K-Means for Four Different Features

68 classifying whether a bounded parking spot was occupied or
69 not occupied by a car in the image.

70 We trained this model on the training dataset (around 250
71 images from the original dataset) and then tested its perfor-
72 mance on the test set (around 90 images). As mentioned
73 earlier, we wrote a function to evaluate the error of the model
74 on the test dataset by calculating the proportion of occupancy
75 classes it predicted wrong, out of all of the examples. The
76 accuracy we got for the baseline RCNN model was 0.737.
77

78 3.2 Updated RCNN Model

79 Overview

80 The updated RCNN model represents an effort to enhance
81 the existing RCNN architecture by introducing an additional
82 fully connected (FC) layer. This adjustment aims to cap-
83 ture more complex features from the region of interest (ROI)
84 warps, with the intention of improving the classification ac-
85 curacy of parking space occupancy.

86 Architecture

87 The enhanced model, RCNN_UPDATED, incorporates an addi-
88 tional FC layer after the backbone ResNet50’s average pool-
89 ing layer. The layer is configured with 1024 neurons, a ReLU
90 activation function, and a dropout of 0.5 for regularization.
91 The output of this layer then feeds into the final classifica-
92 tion layer, which predicts the occupancy status. The backbone’s
93 final FC layer is also adapted to accommodate the reduced
94 feature dimensionality output by the new FC layer.

95 Training Challenges

96 During the training process, we encountered significant com-
97 putational challenges. The Jupyter notebook environment
98 consistently crashed due to excessive RAM usage, likely
99 caused by the increased complexity and size of the model.
100 As a result, we could not complete the training with the newly
101 initialized weights for the additional FC layer.

102 Testing with Random Weights

103 Due to the aforementioned constraints, we proceeded to test
104 the model’s initial behavior by instantiating random weights
105 and biases for the additional FC layer. While these random
106 weights do not carry meaningful learned representations, this

107 step was necessary to verify the updated model’s forward
108 pass functionality. The average percent test error on the train
109 dataset, noted as 0.0736977177718558, indicates the ran-
110 domness of the initial weights and underlines the necessity of
111 proper training. For future work, the RCNN_UPDATED model
112 requires comprehensive training with adequately provisioned
113 computational resources to prevent environment crashes and
114 ensure the model’s convergence. It would also be prudent to
115 explore model optimization techniques to reduce RAM us-
116 age, possibly through model pruning or the implementation
117 of a more efficient training pipeline.

4 Experiments

118 4.1 K-Means

119 We also wanted to do some exploration on different types of
120 features that we could extract from the parking lot datasets.
121 We were able to individually extract color, position, gradient,
122 and edge features and run k-means clustering on them to get
123 assignments and segments, and apply them accordingly to
124 the original image and visualize the results. [See Figures 2
125 and 3].

128 4.2 Perspective Warping

129 What We Tried to Do

130 We looked into using perspective warping to transform high-
131 angle bird’s eye view parking lot images into lower angle
132 views, similar to what you’d see from a camera mounted on
133 a vehicle. The idea was to create a test dataset that closely
134 matched what our parking lot detection and analysis model
135 would encounter in the real world.

136 Our Approaches

137 **Automated point selection** We tried to automatically iden-
138 tify key features and points in the bird’s eye view images and
139 map them to corresponding points in the low angle view. We
140 used techniques like edge detection, corner detection, and line
141 intersections to find parking space boundaries and other ref-
142 erence points that could help guide the transformation.

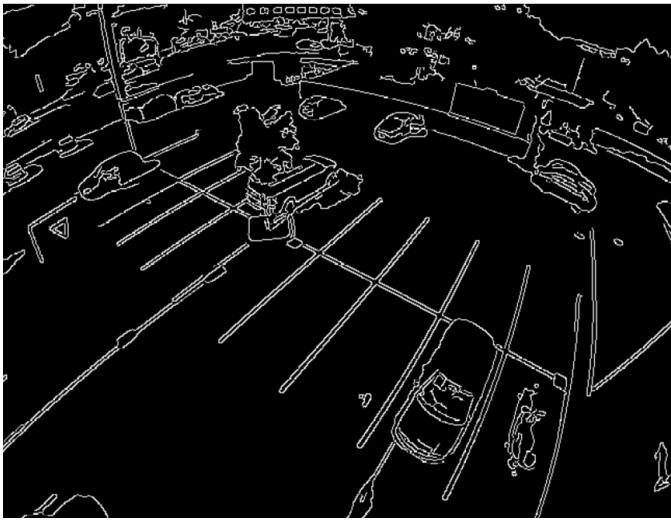


Figure 4: Edge Detection

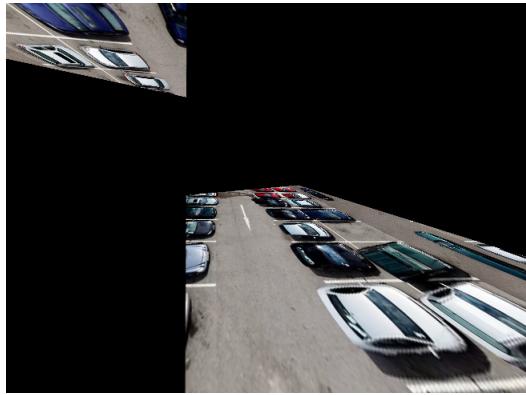


Figure 5: Distorted image when attempting warping



Figure 6: One of the images from the curated street dataset.

143 **Using GitHub dataset annotations** We found a GitHub
 144 dataset [<https://github.com/martin-marek/parking-space-occupancy/tree/main>] that had annotated parking space
 145 coordinates in an annotations.csv file. We thought these
 146 could be used as source points for the perspective transfor-
 147 mation.
 148

149 **Two-step transformation: affine and perspective** Since
 150 the bird’s eye view and low angle view are quite different, we
 151 tried breaking down the transformation into two steps. First,
 152 an affine transformation to adjust the overall orientation and
 153 proportions of the image, and then a perspective transforma-
 154 tion to simulate the depth and foreshortening effects.

155 **Camera model and inverse perspective mapping** In-
 156 spired by MATLAB’s birdsEyeView feature, we tried to
 157 model the perspective warp using camera parameters like focal
 158 length, principal point, camera height, and pitch. We then
 159 applied an inverse perspective mapping to project the bird’s
 160 eye view onto a virtual camera plane.

161 Challenges We Faced

162 **Inconsistent and inaccurate point correspondences** Au-
 163 tomatically detecting matching points between the bird’s eye
 164 and low angle views turned out to be error-prone. The fea-
 165 tures looked quite different and were often occluded in the
 166 low angle view, making it hard to establish reliable point cor-
 167 respondences.

168 **Limitations of annotation data** Using the parking space
 169 coordinates from the GitHub dataset annotations as source
 170 points for the perspective warping didn’t give us good results.
 171 The coordinates didn’t directly correspond to the desired low
 172 angle perspective, so the transformed images ended up dis-
 173 torted and unrealistic.

174 **Oversimplification of 3D to 2D transformation** The two-
 175 step transformation approach (affine followed by perspective)
 176 wasn’t enough to fully represent the complex 3D to 2D map-
 177 ping. The affine step could only approximate the desired
 178 viewpoint, and the perspective warp didn’t fully capture the

nonlinear effects. This led to artifacts like stretched regions
 179 and weird depth discontinuities.
 180

181 **Mismatch between camera model and real-world varia-
 182 tions** Using fixed camera parameters to model the perspec-
 183 tive warp had its drawbacks. In the real world, parking lots
 184 have uneven terrain, obstacles at different heights, and cam-
 185 eras mounted at different positions on vehicles. The simpli-
 186 fying assumptions of our camera model just didn’t hold up
 187 well, resulting in unrealistic distortions.

188 4.3 Creating the street dataset out of the original 189 dataset

190 Given the challenges we faced with perspective warping, we
 191 ultimately decided to take a different approach for creating
 192 our test dataset.

193 We first attempted to find some street images from the in-
 194 ternet and creating a dataset that way, but we realized that the
 195 ROI information was essential input for calculating the scores
 196 during the forward pass.

197 Instead of trying to transform the bird’s eye view images,
 198 we carefully selected a subset of twelve images from the full
 199 training parking lot dataset that most closely resembled the
 200 street parking scenario we were interested in testing.

201 To prepare this abridged test dataset for our model, we cre-
 202 ated a corresponding dataset of regions of interest (ROIs) for

203 each image. These ROIs highlight the specific areas in each
204 image that we wanted our model to focus on during the testing
205 phase [See Figure 6].

206 By using this approach, we were able to create a more realistic
207 and reliable test dataset that closely matched our desired
208 street parking scenario. This allowed us to effectively evaluate
209 our model’s performance without the need for error-prone
210 perspective warping techniques.

211 While this test dataset was smaller in size compared to our
212 original plan, it provided a more accurate representation of
213 the real-world conditions our model would encounter. This
214 strategy enabled us to proceed with testing and validating our
215 parking lot detection and analysis model using data that we
216 could trust.

217 **4.4 Results and Discussion**

218 We were able to run both of our RCNN and RCNN_Updated
219 models on this smaller street dataset and get their accuracies
220 as well. The average test error per image for the RCNN model
221 was 0.0617 and, for the RCNN_Updated file, it was 0.3449.

222 **5 Conclusion and Future Work**

223 Throughout this paper, we have summarized our quantitative
224 results from running the two models on both the original
225 dataset and also the specifically curated street dataset of 12
226 images. Using the original dataset, we saw that the testing
227 error had improved slightly with the updated RCNN model
228 but was mostly the same between the two models. Interestingly,
229 the RCNN model did not perform significantly better
230 despite adding another layer to it. However, because we had
231 to test with random weights, this could have negatively impacted
232 the performance of the updated model, and in the future,
233 we plan to improve upon this. Finally, we found that
234 the original model still performed well on the street dataset
235 that we created, with the RCNN updated model performing
236 much worse, however. Despite this, we were able to target
237 our RCNN model towards the street detection problem and
238 identified one area of improvement, which was to create a
239 dataset and model such that the model would not be dependent
240 on ROIs data to make predictions.

241 Ultimately, we found that adding the additional complexity
242 did not necessarily improve performance in certain scenarios.
243 By performing perspective warping and k-means, we encountered
244 a few challenges and visualizations that will be useful
245 for future work.

246 In the future, we plan to create models that do not depend on
247 ROI information for testing, as well as seeing how we can apply
248 this work to AutoPark from Tesla. Recent developments
249 include advancing AutoPark capabilities without reliance on
250 an adjacent car requirement, and we hope to help enhance the
251 solution to this problem with UrbanPark.

252 **6 Contributions**

253 Cynthia wrote the k-means visualizations of the images,
254 created the street parking dataset and wrote the percent error
255 evaluation function. Abhi wrote the RCNN_Updated model
256 and the code and visualizations for perspective warping.
257 Both group members set up the data loading and baseline

RCNN model and contributed equally in writing the report. 258
259

260 **7 Appendix**

261 The code for our project is linked here: 262
262 <https://drive.google.com/drive/folders/1DU3zl0NQvP7K7cpRJr9VVMZB9nfqJYTV?usp=sharing> 263
264

265 **References**

- 266 [Acharya *et al.*, 2018] Debaditya Acharya, Weilin Yan, and 267
267 Kourosh Khoshelham. Real-time image-based parking oc-
268 cupancy detection using deep learning. 04 2018. 269
- 269 [Amato *et al.*, 2017] Giuseppe Amato, Fabio Carrara, Fab- 270
270 rizio Falchi, Claudio Gennaro, Carlo Meghini, and Clau- 271
271 dio Vairo. Deep learning for decentralized parking lot 272
272 occupancy detection. *Expert Systems with Applications*, 273
273 72:327–334, 2017. 274
- 274 [Anuj Khandelwal, 2021] Anuj Khandelwal. Bird’s eye view 275
275 transformation. <https://gist.github.com/anujonthemove/7b35b7c1e05f01dd11d74d94784c1e58>, 2021. Accessed: 276
276 2023-MM-DD. 277
- 278 [Marek, 2021] Martina Marek. Image-based parking space 279
279 occupancy classification: Dataset and baseline. *ArXiv*, 280
280 [abs/2107.12207](https://arxiv.org/abs/2107.12207), 2021. 281
- 281 [Matsuda *et al.*, 2021] Akihiro Matsuda, Tomokazu Matsui, 282
282 Yuki Matsuda, Hirohiko Suwa, and Keiichi Yasumoto. A 283
283 system for real-time on-street parking detection and vi- 284
284 sualization on an edge device. In *2021 IEEE Interna- 285
285 tional Conference on Pervasive Computing and Commu- 286
286 nications Workshops and other Affiliated Events (PerCom 287
287 Workshops)*, pages 227–232, 2021. 288